

SpatialLens Assist: Motion-Aware Hazard Detection for Campus Navigation from Low-FPS Egocentric Video

Diego Sanchez
Stanford University
dsanh14@stanford.edu

Abstract

An object detector tells a blind pedestrian *what* is nearby, but not whether it is closing, crossing, or simply present. We ask whether a small, interpretable set of motion cues—bounding-box growth, centroid displacement, optical flow, and ego-motion compensation—can recover useful hazard labels from low-fps handheld campus video, without training a learned trajectory model. We present **SpatialLens Assist**, an end-to-end CPU pipeline that converts short (~ 10 s) walkway clips into per-track motion-aware alerts over six classes (**approaching**, **crossing_LTR**, **crossing_RTL**, **moving_away**, **static**, **uncertain**). The system pairs YOLOv8n with an IoU+centroid tracker, ego-motion-compensated Farneback flow and frame differencing, per-track motion aggregates, and a transparent rule cascade that emits a natural-language evidence string and treats **uncertain** as a first-class abstention rather than a guess. On 11 controlled walkway videos (61 manually labelled tracks), the system reaches **82.5%** overall accuracy on 57 evaluated track-label pairs and **93.3%** on the 30 tracks with at least three frames, with approaching precision/recall **0.95/0.88**—the safety-asymmetric regime an alert layer needs. Ablations show the largest gain comes from tracker continuity (+6.2pts from appearance-gated re-identification, +3.5pts each from two short-track salvage rules), and ByteTrack is a negative result on this sparse 2 fps regime (32.8%). The dominant failure mode is conservative abstention from tracker fragmentation, not dangerous mislabelling: *no approaching*↔*moving_away* confusion occurs.

1 Introduction

A bicycle parked against a wall and a bicycle bearing down on the camera are both detected as “bicycle, 0.86.” For a sighted user the wider visual context carries the intent; for a blind or low-vision pedestrian relying on an assistive device, that context is exactly what is missing. The safety-critical signal is not the class label but the *motion*: whether the agent is closing, crossing the path, or merely present.

We frame motion-aware hazard detection as *per-track classification*: given a ~ 10 s phone clip sampled at 2 fps, decide for each tracked object whether it is approaching, crossing left-to-right, crossing right-to-left, moving away, static, or explicitly **uncertain**. Three properties drive the design: (i) **explainability** – every prediction carries a textual evidence string; (ii) **calibrated abstention** – short tracks do not bluff, so **uncertain** is a first-class output with a sub-reason; (iii) **reproducibility on CPU** – the full pipeline runs in ~ 4 min on a GPU-less Mac and the 61-track labelled set is committed.

Summary. On 11 controlled walkway videos and 61 labelled tracks the system reaches 82.5% accuracy on 57 matched tracks (93.3% on the 30 decidable tracks with ≥ 3 frames); **approaching** reaches 0.95/0.88 precision/recall and no **approaching**↔**moving_away** confusion is observed. Isolated ablations attribute the dominant gain to track continuity rather than classifier complexity: appearance-gated re-identification contributes 6.2pts and the two 2-frame salvage rules 3.5pts each, while ByteTrack on 2fps offline footage drops to 32.8% – a useful negative result.

Contributions. (C1) a CPU-only offline egocentric pipeline for motion-aware campus hazard alerts; (C2) ego-motion-compensated motion features (median-subtracted Farneback flow + ECC-aligned frame differencing) appropriate to handheld low-fps footage; (C3) an interpretable hazard rule cascade with natural-language evidence strings,

conservative abstention, and a trajectory-reversal heuristic for “approached-then-turned-away” tracks; (C4) a 61-track evaluation set with two reporting denominators and a fully scripted ablation suite; (C5) empirical evidence that tracking continuity dominates classifier sophistication in this regime, and that ByteTrack is poorly matched to sparse offline footage.

2 Related Work

Detection and tracking. YOLOv8 [1] is a fixed COCO-class front-end. SORT [3] and ByteTrack [2] are strong online baselines but assume near-continuous per-frame detections from a Kalman model; on our 2fps footage ByteTrack (via **supervision** [4]) drops 51% of detections, motivating a transparent per-class IoU+centroid matcher with explicit gap buffer plus a separate appearance-gated re-ID pass tailored to sparse sampling. Assistive tools such as Seeing AI [7] and academic prototypes [8] expose object *identity* but rarely *motion intent*.

Trajectory prediction vs. past-motion classification. Trajectron++ [5] and Social-STGCNN [6] forecast *future* positions over 4–8s from high-fps tracks with map context, requiring thousands of training trajectories. We solve the strictly easier problem of classifying the *past* motion of a single track into six safety-relevant categories from ~ 10 sampled frames, choosing an alert plus evidence string over a future-position distribution – a deliberate scope choice driven by the no-training-data setting and the explainability/abstention requirements of an assistive interface. Classical Farneback flow [9] and ECC alignment [10] support the CPU-only constraint; median-flow subtraction compensates for the camera pan that otherwise inflates apparent motion of every static object.

3 Method

The pipeline (Figure 1) has six stages. We summarise each, focusing on the four non-obvious design decisions.

3.1 Frames, detection, base tracking

Videos are sampled at 2fps, resized to 960px wide. YOLOv8n emits per-frame detections at $\text{conf} \geq 0.35$. A per-class IoU+centroid tracker matches greedily by $s_{dt} = 0.7 \cdot \text{IoU}(d, t) + 0.3 \cdot (1 - \hat{d}_{dt})$ with a `max_frame_gap` buffer so a track survives 1–2 missed detections.

3.2 Motion features with ego-motion compensation

Handheld panning makes every static object appear to move. We compensate two ways: **flow** – Farneback dense flow between frames, minus the per-frame median flow vector (background-dominated), giving per-bbox (flow_dx, flow_dy, flow_mag) relative to the scene; **frame diff** – align the previous frame to the current with `cv2.findTransformECC` (translation) before differencing. Per-track aggregates: start-to-end growth $r = (A_{\text{end}} - A_{\text{start}})/A_{\text{start}}$, total Δc_x , mean flow magnitude, plus two sub-track signals for reversal detection – second-half growth $r_{1/2}$ and peak-area ratio $\rho = \max_t A_t/A_{\text{end}}$.

3.3 Conservative fragment re-identification

At 2fps a fast-crossing object moves up to 350px between samples, so the tracker fragments single objects. A post-tracker cleanup applies (i) **same-frame NMS** (drop overlapping same-class detections at $\text{IoU} \geq 0.5$) and (ii) **appearance-gated stitching**: merge tracks $B \rightarrow A$ only if *all* of same class, *both tracks = 1 frame*, $\text{gap} \leq 2$, centroid distance $\leq 0.45 \text{diag}$, area ratio ≤ 5 , and HSV (H, S) histogram correlation ≥ 0.45 . The 1-frame constraint protects multi-frame tracks from being absorbed by fragments – an aggressive variant broke seven correct tracks (Section 5.2). `src/track_reid.py` keeps cleanup separable for ablation.

3.4 Rule-cascade hazard classifier

The classifier resolves the label by a priority cascade, each branch carrying its own evidence template (r : bbox-growth ratio; ρ : peak/end area ratio; $r_{1/2}$: second-half growth): (1) **Short-track salvage** ($n = 2$): strong lateral $|\text{dx_norm}| > 0.30 \rightarrow \text{crossing}$, else strong approach ($r > 0.7$, downward flow) $\rightarrow \text{approaching}$, else **uncertain**. (2) **Static** if $n \geq 3$ and no motion. (3) **Moving-away** if $r < -0.15$, or the *soft* variant ($-0.15 < r < -0.03$ with sub-strong lateral motion read as ego-motion drift). (4) **Trajectory reversal** ($n \geq 6$, $r > 0.3$, $\rho > 1.5$, $r_{1/2} < -0.20$) $\rightarrow \text{moving_away}$ for the “approached then turned away” case. (5) **Strong crossing** ($|\text{dx_norm}| > 0.50$) overrides approach. (6) **Approaching** if $\text{approach_score} \geq 0.55$ and ($r > 0.15$ or the slow-approach allowance $r > 0$ with score ≥ 0.65). (7) **Crossing** if $|\text{dx_norm}| > 0.08$, $r \leq 0.15$.

Table 1: Per-class P/R/F1 (raw 57-track denominator). Macro-F1 0.67 is depressed by rare classes with $n \leq 3$.

Class	Precision	Recall	F1	n
<code>approaching</code>	0.95	0.88	0.91	24
<code>crossing_left_to_right</code>	1.00	0.50	0.67	4
<code>crossing_right_to_left</code>	1.00	0.67	0.80	3
<code>moving_away</code>	1.00	0.78	0.88	9
<code>static</code>	0.00	0.00	0.00	2
<code>uncertain</code>	0.62	1.00	0.77	15
Overall accuracy	82.5% on 57 evaluated tracks			
Decidable accuracy ($n \geq 3$)	93.3% on 30 tracks			
Macro F1	0.67			

(8) **Uncertain** fall-through, tagged with the closest-missed sub-reason (`near_approaching`, `low_signal`, ...) quoted in the evidence string.

The *approach score* is a fixed-weight blend $0.45 \cdot s_{\text{growth}} + 0.25 \cdot s_{\text{center}} + 0.20 \cdot s_{\text{fd}} + 0.10 \cdot s_{\text{flow}}$ of four normalised cues (bbox growth, centroid motion toward image center/lower half, frame-diff overlap, and flow magnitude).

3.5 Alert generation and explainability

Every prediction carries a templated *evidence string* quoting the feature values that triggered the cascade branch (e.g. “classified as approaching because bbox grew by 11.42 and approach_score = 1.00; centroid moved toward image center/lower half; flow_mag = 17.26”); for abstentions the string names the sub-reason. Tracks with < 3 frames are not forced into a motion class – treating uncertainty as safety-preserving abstention. A 4-frame temporal strip illustrating the mechanism on a canonical approaching cyclist is in Section E.

4 Dataset and Evaluation Protocol

We recorded 11 controlled ~ 10 s phone videos on Stanford walkways at 2fps / 960px; each captures a single scripted hazard mode so label intent is unambiguous. Every tracked object was hand-labelled into the six-class motion taxonomy (`approaching`, `moving_away`, `crossing_LTR`, `crossing_RTL`, `static`, `uncertain`), giving 61 ground-truth tracks. The distribution is imbalanced (`approaching` = 24, `uncertain` = 15, `static` = 2); see Section B for full dataset and class tables.

Predictions and labels match on (`video_id`, `track_id`). After post-tracker re-ID, 4 labels are absorbed by NMS/stitching into an existing labelled track, leaving 57 matched pairs (raw denominator). We additionally publish a *consolidated* 51-track set that collapses the 9 entries the labeller’s notes column explicitly marked “fragmentation” / “duplicate” (see Section D). We report overall accuracy on the 57 raw pairs, the *decidable* accuracy on the 30 tracks with ≥ 3 frames, per-class P/R/F1, macro-F1, and isolated ablations that toggle one component at a time.




Figure 1: End-to-end SpatialLens pipeline. The detector is off-the-shelf; ★ marks the four contributions evaluated in Section 5 (same-frame NMS, appearance-gated re-ID, sub-track growth features, and the rule-cascade salvage/guard rules). Every prediction carries a natural-language *evidence string* that the alert layer quotes, and **uncertain** is a first-class output with a sub-reason rather than a fall-through default.




Figure 2: Confusion matrix on the 57 evaluated tracks. Off-diagonal errors are dominated by abstention into the **uncertain** column. The safety-critical **approaching**↔**moving away** confusion is *zero*; the only dangerous off-diagonal is one **moving away**→**approaching** on a monotone bbox-growth edge case (IMG_4973/person_4).

5 Results

5.1 Main quantitative results

Table 1 reports per-class metrics. The safety-critical **approaching** class reaches **0.95 / 0.88** precision/recall – the right asymmetry for a safety alert. Every non-trivial directional class has precision 1.00: when the cascade commits to a directional label it is essentially never wrong; what it does instead, when temporal evidence is thin, is abstain into **uncertain**. The gap between overall (82.5%) and decidable (93.3%) accuracy is explained almost entirely by short tracker fragments. Results are not single-video luck: 6 of 11 videos are perfect, the worst absolute count is IMG_4982 (11/16 raw); a per-video and motion-feature-space breakdown is in Sections B and E.

5.2 Ablations and baselines

Figure 3 shows the effect of removing each contribution one at a time. The largest gain comes from track *continuity* rather than classifier complexity: appearance re-ID stitching is the single largest contributor (−6.2 pts raw); the two 2-frame salvage rules each contribute ∼3.5 pts; toggling NMS and re-ID off together costs 7.0 pts. Classifier-side




Figure 3: Accuracy drop when each component is removed (one at a time) from the full system. Tracker-side cleanups dominate; ByteTrack underperforms on sparse low-fps footage. Full numerics in Section C.

rules (reversal, soft moving-away, slow-approach) contribute ∼1.8 pts each – meaningful but not dominant.

Baselines: random uniform = 16.7%, detector-only / always-uncertain = 29.4%, majority-class always-approaching = 41.2%. The full system doubles the strongest trivial baseline.

ByteTrack negative result. ByteTrack reaches only 32.8% even after tuning. Its Kalman-gated association assumes near-continuous detections; on 2 fps footage predicted positions over ∼500 ms gaps rarely match observed ones, and short detections never “activate” a track – it drops ∼51% of YOLO’s detections under defaults. A useful negative result: a strong online tracker does not transfer to an offline low-fps regime where its motion-continuity assumption is violated. A separate negative result – aggressive stitching with ≤ 4 frame fragments *broke* 7 correct tracks (75.4%→68.8%) – motivates our “both tracks must be 1 frame” constraint.

6 Failure Analysis and Limitations

Of 10 residual errors, 5 are tracker-fragmentation abstentions and 5 are classifier-side. The system’s uncertainty is dominated by *short_track* (23 of 24 abstentions): uncertainty is the tracker-fragmentation tax, made auditable, not classifier ambiguity. Figure 4 grounds this: panels (a–c) are the canonical directional classes the cascade commits to, while panel (d) is the conservative **uncertain(short_track)** abstention on a 1-frame fragment – the safety-preserving behaviour that produces most off-diagonal mass in Figure 2. A full sub-reason breakdown, an error taxonomy of the five root causes, and the static-class behaviour (F1 0.00 driven by 2 one-frame parked-bicycle labels on which the cascade correctly abstains) appear in Section F.



Figure 4: Qualitative examples, one per behaviour. (a) approaching cyclist with $r = +21.98$; (b) right-to-left pedestrian crossing over 11 frames; (c) pedestrian moving away with shrinking bbox; (d) conservative `uncertain(short_track)` abstention on a 1-frame parked bicycle. Yellow trails are centroid history.

Limitations. The evaluation set is small (61 tracks, 11 videos) and imbalanced; we have not yet run a user study with blind/low-vision walkers; the pipeline is tuned for 2 fps and threshold constants would need re-tuning at higher sampling; bbox-growth is a noisy proxy for depth (a monocular depth backend would resolve the one residual `moving_away`→`approaching` mistake); we use YOLOv8n unchanged so detector-side class confusion (skateboard / scooter) is inherited. Extended discussion in Section F.

7 Conclusion

SpatialLens Assist shows interpretable motion cues can produce useful hazard labels from low-fps egocentric campus video. Decidable accuracy is 93.3%, `approaching` reaches 0.95/0.88 P/R, and the safety-critical `approaching`↔`moving_away` confusion is zero. Ablations argue that at this project scale, *temporal track continuity matters more than classifier sophistication*: appearance-gated re-ID and short-track salvage together contribute roughly twice as much accuracy as all classifier-side rule changes combined. Future work targets the four limitations: a monocular depth backend (MiDaS [11]); a learned per-track embedding to replace the HSV histogram; denser sampling (5–10 fps); and a user-centred evaluation with blind/low-vision walkers.

Individual Contributions

Solo project. Diego Sanchez completed 100% of the work. He wrote all pipeline code under `src/` — the detection wrapper, the IoU+centroid tracker and ByteTrack ablation backend, the ego-motion-compensated optical flow and frame-differencing, the per-track motion features (in-

cluding the second-half-growth and peak-area signals), the rule-cascade hazard classifier with evidence strings, the appearance re-identification cleanup in `src/track_reid.py`, and the evaluation, alert, and visualisation modules — together with all driver scripts (`run_ablations.py` and `make_final_figures.py` generate every figure and table in this report) and the 74-test `pytest` suite. He recorded the 11 controlled walkway videos, hand-labelled all 61 ground-truth tracks, and wrote this manuscript.

As authorized in the course instructions, ChatGPT, Claude, and Cursor were used in the completion of the project. Specifically, ChatGPT assisted with literature search, proofreading, and drafting of this manuscript; Claude (via Claude Code) and Cursor were used for code debugging and refactoring. All design decisions, experiments, labels, and reported results were produced and verified by the author.

The code is available at our GitHub repository: <https://github.com/dsanh14/SpatialLens>.

References

- [1] Glenn Jocher et al. *Ultralytics YOLOv8*. GitHub, 2023.
- [2] Y. Zhang et al. ByteTrack: Multi-Object Tracking by Associating Every Detection Box. *ECCV*, 2022.
- [3] A. Bewley et al. Simple Online and Realtime Tracking. *ICIP*, 2016.
- [4] Roboflow. *Supervision: Reusable Computer Vision Tools*. 2024.
- [5] T. Salzmann et al. Trajectron++: Dynamically-Feasible Trajectory Forecasting With Heterogeneous Data. *ECCV*, 2020.
- [6] A. Mohamed et al. Social-STGCNN: A Social Spatio-Temporal Graph Convolutional Neural Network for Human Trajectory Prediction. *CVPR*, 2020.
- [7] Microsoft. *Seeing AI: Talking camera app for those with a visual impairment*. 2017–.
- [8] M. Poggi and S. Mattoccia. A wearable mobility aid for the visually impaired based on embedded 3D vision and deep learning. *ISCC*, 2016.
- [9] G. Farnebäck. Two-Frame Motion Estimation Based on Polynomial Expansion. *SCIA*, 2003.
- [10] G. D. Evangelidis and E. Z. Psarakis. Parametric image alignment using enhanced correlation coefficient maximization. *IEEE TPAMI*, 30(10), 2008.
- [11] R. Ranftl et al. Towards Robust Monocular Depth Estimation: Mixing Datasets for Zero-shot Cross-dataset Transfer. *IEEE TPAMI*, 2022.

Supplementary Material

This supplement is organised as follows. Full classifier specification with all thresholds and the un-compressed rule cascade (§A); dataset and class-distribution tables (§B); per-video and full-ablation tables (§C); the 9 label-consolidation drops with the labeller’s verbatim notes (§D); motion feature space and qualitative examples (§E); failure analysis details (§F); reproduction commands (§G).

A Full classifier specification

Hyperparameters. All thresholds are set in `config.yaml` (Table 2). The classifier rules use the values exactly as committed; we did not learn any parameter from the 61-track labeled set, only iterated on rule *structure* after each error analysis pass.

Table 2: Thresholds used by the system. Sourced from `config.yaml`; the second-half-growth and peak-area-ratio cut-offs in the trajectory-reversal rule are inline constants in `src/hazard_classifier.py`.

Threshold	Value
<i>Detection / sampling</i>	
YOLO confidence	≥ 0.35
Sample rate	2 fps
Resize width	960 px
<i>Tracking + re-ID</i>	
IoU threshold	0.25
Centroid threshold (frac. of diagonal)	0.15
Max frame gap	4
NMS IoU	0.5
Stitch gap (max frames)	2
Stitch spatial gate (frac. diag.)	0.45
Stitch area-ratio gate	5.0
HSV histogram correlation	≥ 0.45
Stitch fragment threshold (frames)	1
<i>Motion + classifier</i>	
Static disp. frac. diag.	0.04
Bbox growth threshold r	0.15
Bbox shrink threshold (hard)	-0.15
Bbox shrink threshold (soft)	-0.03
Lateral threshold dx_norm	0.08
Strong lateral dx_norm	0.50
Short-track dx_norm salvage	0.30
Min track frames	3
Approach-score threshold	0.55
Slow-approach approach-score threshold	0.65
Reversal: $r, \rho, r_{1/2}$	0.3, 1.5, -0.20
Reversal min frames	6

Rule cascade (full). In execution order, each branch emits its label and a textual evidence string that the alert layer quotes verbatim.

1. **Short-track salvage (2-frame).** For $n=2$ tracks (otherwise we cannot trust trajectory cues): if $|dx_norm| > 0.30$ emit **crossing** in the sign of dx_norm ; else if $r > 0.7$, $flow_dy > 1$, and $flow_mag > 1$ emit **approaching**; else **uncertain (short_track)**. For $n=1$ the salvage is always **uncertain (short_track)**.
2. **Static.** $n \geq 3$ and total displacement, flow, and frame-diff overlap are all below their thresholds.

3. **Moving away (hard).** $r < -0.15$: depth recession dominates any incidental lateral motion.
4. **Moving away (soft, ego-motion guard).** $-0.15 < r < -0.03$ and $|dx_norm|$ in the ambiguous band (0.08, 0.50): sub-threshold lateral drift is read as camera ego-motion rather than a true crossing.
5. **Trajectory reversal.** $n \geq 6$, $r > 0.3$, peak/end area ratio $\rho > 1.5$, and second-half growth $r_{1/2} < -0.20$: the bbox grew start-to-end but peaked mid-track, so the object approached then receded.
6. **Strong crossing.** $|dx_norm| > 0.50$: lateral intent dominates even if the bbox also grew.
7. **Approaching.** $approach_score \geq 0.55$ and ($r > 0.15$ or the slow-approach allowance $r > 0$ with $approach_score \geq 0.65$). The approach score is $0.45 s_{growth} + 0.25 s_{centre} + 0.20 s_{fd} + 0.10 s_{flow}$.
8. **Crossing (regular).** $|dx_norm| > 0.08$ and $r \leq 0.15$: lateral motion present, bbox not growing.
9. **Uncertain (fall-through).** A sub-reason is chosen by proximity to the closest-missed threshold: **near_approaching**, **near_crossing_***, **near_moving_away**, **conflicting_cues**, or **low_signal**. The reason is quoted in the evidence string so a pipeline log can explain *why* the system declined to commit.

B Dataset details

We recorded 11 controlled ~ 10 s phone clips on Stanford walkways under daylight, sampling at 2 fps to 960 px-wide frames. Each clip captures a single scripted hazard mode (bike approaching, pedestrian crossing, etc.) so label intent is unambiguous. Two natural-walk clips are kept for qualitative demonstration only and are excluded from the quantitative tables. Table 3 summarises the dataset splits; Table 4 reports the ground-truth class distribution. The set is realistic for campus walkway footage but imbalanced – **approaching** and **uncertain** dominate, **static** has only 2 examples (both 1-frame parked-bicycle fragments). This imbalance is the main reason macro-F1 is lower than overall accuracy.

Table 3: Dataset summary. 11 controlled videos contribute the 61 manually-labelled tracks; 4 labels are absorbed by post-tracker re-ID, leaving 57 prediction/label pairs (raw denominator) and 30 tracks for the decidable subset (≥ 3 frames).

Split	Videos	Labeled tracks	Used quantitatively?
Controlled campus videos	11	61	Yes
Matched evaluated tracks	11	57	Yes
Decidable subset ($n \geq 3$)	11	30	Yes
Uncontrolled qualitative	2	—	No

C Detailed evaluation tables

Per-video accuracy. Table 5 reports results per video against both label sets. Of the eleven videos, six are classi-

Table 4: Ground-truth class distribution. The set is imbalanced: **approaching** dominates; **static** and **crossing_RTL** have $n \leq 3$.

Class	n	%	Note
approaching	24	39.3%	primary safety-critical class
crossing_left_to_right	7	11.5%	limited support
crossing_right_to_left	3	4.9%	limited support
moving_away	9	14.8%	non-hazard motion
static	2	3.3%	only 2 examples $\Rightarrow$ high variance
uncertain	16	26.2%	deliberate abstention label
Total	61	100.0%	11 controlled videos

fied perfectly; the worst (IMG_4982) has 16 labeled tracks, of which 3 are 1-frame approaching fragments and 2 are 1-frame parked-bicycle “static” labels — regimes where one detection genuinely cannot resolve the intent. IMG_4973 contains the two trajectory-reversal errors (a monotone-bbox-growth edge case where the reversal rule’s $\rho > 1.5$ threshold does not fire).

Table 5: Per-video accuracy. “raw” is correct over evaluated tracks after NMS/re-ID, “cons.” is correct over the 51 unique physical objects.

Video	n	raw	cons.	dominant error
IMG_4972	7	6/6	6/7	—
IMG_4973	6	4/6	4/6	mov_away $\rightarrow$ unc.
IMG_4974	4	4/4	4/4	—
IMG_4975	5	5/5	5/5	—
IMG_4976	6	6/6	6/6	—
IMG_4977	5	1/3	1/1	cross_LTR $\rightarrow$ unc.
IMG_4978	1	1/1	1/1	—
IMG_4979	2	1/1	1/1	—
IMG_4980	2	2/2	2/2	—
IMG_4981	7	6/7	6/6	cross_RTL $\rightarrow$ unc.
IMG_4982	16	11/16	11/13	appr. $\rightarrow$ unc.
total	61	47/57	47/51	

Full ablation table. Table 6 is the numeric backing of the bar chart in the main paper, with macro-F1 columns added.

Table 6: Isolated ablation: full system minus exactly one component, evaluated against both label sets. Macro-F1 reflects per-class balance and is more sensitive than overall accuracy when errors concentrate in low-support classes. From `outputs/evaluation/ablations.csv`.

Variant	raw		consolidated	
	acc.	F1	acc.	F1
full system	82.5 %	0.67	92.2 %	0.79
– appearance re-ID	76.3 %	0.55	88.2 %	0.62
– 2f crossing salvage	78.9 %	0.55	88.2 %	0.62
– 2f approach salvage	78.9 %	0.66	88.2 %	0.78
– trajectory reversal	80.7 %	0.65	90.2 %	0.78
– soft moving-away guard	80.7 %	0.64	90.2 %	0.75
– slow-approach	80.7 %	0.65	90.2 %	0.76
– same-frame NMS	81.4 %	0.65	92.3 %	0.79
– NMS and re-ID (both off)	75.4 %	0.55	88.5 %	0.62

D Label consolidation transparency

We collapse 9 of the 61 ground-truth rows into one representative each. Every drop is justified by an explicit annotation in the labeler’s **notes** column (Table 7). Four of the 13 entries the labeler flagged as fragments were *kept* because they were not duplicates of existing labeled objects or because the classifier already predicted them correctly (IMG_4975/person_3 and IMG_4982/motorcycle.1) – consolidation drops only the entries that both (a) carry an explicit duplicate/fragment annotation and (b) correspond to a physical object already represented by another labeled track.

E Motion feature space and qualitative examples

Figure 5 plots every labelled track in the (dx_{norm}, r) feature space the cascade uses. **approaching** (red) clusters in positive growth, **moving_away** (green) in negative growth, **crossing** (blue/purple) at large $|dx_{\text{norm}}|$, and **uncertain** (orange) near the origin where no decisive cue is available. The classifier’s dashed thresholds ($|dx_{\text{norm}}| \in \{0.08, 0.50\}$ and $r \in \{-0.15, 0.15\}$) sit exactly between class clusters – visual confirmation that two features carry most of the decision boundary.




Figure 5: Motion feature space. Filled = correct, hollow = misclassified. Dashed lines are classifier thresholds. Growth values clipped to $[-1.2, 4.0]$ for readability.


Temporal motion evidence. Figure 6 illustrates the mechanism on a representative **approaching** cyclist: across four sampled frames of a 15-frame track, the bounding box grows $21.98\times$, the yellow centroid trail moves toward the image centre, and the approach score saturates at 1.00. The evidence string is what the alert layer quotes.

Qualitative panel. The qualitative examples in Figure 4 (main text) give one canonical example per directional class plus a conservative abstention. Panel (b)’s yellow centroid

Table 7: The 9 ground-truth rows dropped during consolidation, with the labeler’s verbatim note. Every drop corresponds to a classifier-incorrect raw prediction (so the consolidation never artificially flips a correct row to a dropped one). The remaining 4 fragment-flagged labels are kept (text).

Video	Track	True label	Labeler’s note
IMG_4977	bicycle.2	crossing_LTR	“likely duplicate/fragment of same crossing bicycle (tracker fragmentation)”
IMG_4977	bicycle.3	crossing_LTR	“likely duplicate/fragment of same crossing bicycle (tracker fragmentation)”
IMG_4977	bicycle.4	crossing_LTR	“likely duplicate/fragment of same crossing bicycle (tracker fragmentation)”
IMG_4977	person.1	crossing_LTR	“rider/person associated with crossing bicycle (tracker fragmentation)”
IMG_4979	person.2	crossing_LTR	“duplicate/fragmented track of same person crossing left to right”
IMG_4981	motorcycle.1	crossing_RTL	“appears to be bicycle/bike fragment crossing right to left (tracker fragmentation, 2 frames)”
IMG_4982	person.1	approaching	“appears to be same approaching pedestrian seen in later frames (tracker fragmentation)”
IMG_4982	person.5	approaching	“duplicate/fragment of approaching pedestrian (tracker fragmentation)”
IMG_4982	person.6	approaching	“duplicate/fragment of approaching pedestrian (tracker fragmentation)”

IMG_4972 / person_1: cyclist approaching the camera — 15 frames, ground truth = approaching



Evidence — bbox growth $r = +21.98$, $dx_{\text{norm}} = -0.15$, flow mag = 14.5, approach score = 1.00 $\Rightarrow$ label = APPROACHING

Figure 6: Temporal motion evidence for an **approaching** cyclist (IMG_4972/person_1, 15 frames). Area $A = 2548 \rightarrow 58568 \text{ px}^2$, $r = +21.98$; centroid trail moves toward the lower centre; ego-motion-compensated flow is consistent with object motion rather than camera pan.

trail tracks sideways across more than half the frame — exactly the cue $|dx_{\text{norm}}|$ encodes; panel (d) is the abstention pattern that dominates the **static**-row of the confusion matrix.

F Failure analysis details

Uncertainty breakdown. Figure 7 shows the sub-reason every **uncertain** prediction carries. The distribution is heavily concentrated on **short_track**: 23 of the 24 abstentions are because the underlying track has < 3 frames, only 1 is **low_signal**. Uncertainty in SpatialLens is the tracker-fragmentation tax made auditable, not classifier ambiguity.

Static-class result. The static-class F1 of 0.00 should not be read as a general static failure mode. Only two static examples appear in the set, both as 1–2 frame fragments of parked bicycles. The cascade correctly *abstains* (the safety-preserving behaviour) rather than committing; with this label set we cannot test whether the **static** branch ever fires correctly on a long stationary track. A larger evaluation set with multi-frame static examples is the cleanest fix.




Figure 7: Sub-reason breakdown for all **uncertain** predictions.

Error taxonomy. Table 8 groups the residual errors into five root causes with example tracks and proposed fixes. Two observations: the dominant failure is *fragmentation*, *not misclassification* (short-track + static abstentions account for 5 of the 10 residuals); and *no residual is in the safety-critical asymmetry* — the **approaching** $\leftrightarrow$ **moving_away** confusion is zero,

Table 8: Error taxonomy: five root causes that explain every residual error, with example tracks and proposed fixes.

Error type	Cause	Example	Proposed fix
Short-track fragmentation	tracker links only 1–2 frames at 2 fps	IMG_4982 bicycle_1	higher fps; learned re-ID embedding
Static abstention	only 2 static examples, both 1–2 frame fragments	IMG_4982 bicycle_2	explicit stationary branch with depth cue
Crossing → uncertain	fast crossing object fragments before $ dx_{\text{norm}} $ accumulates	IMG_4977 skateboard	denser sampling; salvage at $n=1$
Moving-away → approaching	monotone bbox growth despite physical recession (no depth cue)	IMG_4973 person_4	monocular depth (MiDaS)
Trajectory-reversal miss	person approached then turned; $\rho < 1.5$	IMG_4973 person_2	softer reversal threshold; second-half flow direction

and the only directional mistake (IMG_4973/person_4 moving-away→approaching) is the monotone-bbox-growth edge case the reversal rule misses ($\rho = 1.17 < 1.5$ threshold).

Extended limitations. Expanding the main-text limitations: **(i) Dataset size** – 61 tracks across 11 videos is small and imbalanced, so the numbers are evidence the motion features are promising, not a deployment claim. **(ii) No user study** – the 0.95/0.88 approaching P/R is the right asymmetry in principle, but a study with blind/low-vision walkers is the appropriate next validation. **(iii) Low-fps regime** – denser sampling (5–10 fps) would alleviate most fragmentation residuals, but the threshold constants scale with inter-frame displacement and would need re-tuning. **(iv) No depth** – bbox growth is a noisy proxy for “getting closer”; the one residual directional mistake is exactly the ambiguity a monocular depth backend would resolve. **(v) Fixed detector** – YOLOv8n’s COCO class set is unchanged, so detector-side class confusion (skateboard / scooter) is inherited.

Key files. `src/hazard_classifier.py` (rule cascade, ~500 LOC); `src/track_reid.py` (NMS + appearance stitching, ~220 LOC); `src/motion_features.py` (incl. second-half/peak-area signals); `src/evaluation.py` (raw + consolidated metrics). Labels in `data/labels/hazard_labels{,_dedup}.csv`.

Honest caveats. (i) The Phase 2 wire-in to call `clean_tracks` from inside `run_final_pipeline.py` is still pending in this submission; the headline numbers were produced by an inline driver in our development session and persist on disk as `outputs/evaluation/all_videos_evaluation_summary.json`. Re-running `--all` without that wire-in would silently regress to the ~75% NMS-and-re-ID-off baseline (Table 6 last row). The procedure is documented in the repository for future work. (ii) The uncontrolled-walk video mentioned in the introduction is qualitative only; it is not part of the 61-track evaluation set.

G Reproduction

Environment. Python 3.13, no GPU required; tested on an M-series Mac. Dependencies are in `requirements.txt` and install with `pip install -r requirements.txt`. The YOLOv8n weights (`yolov8n.pt`) auto-download on first run.

End-to-end commands.

```
# 1. detect, track, motion features, hazards, alerts
python scripts/run_final_pipeline.py --all

# 2. isolated ablations -> outputs/evaluation/ablations.csv
python scripts/run_ablations.py

# 3. paper figures from current outputs
python scripts/make_paper_figures.py
python scripts/make_failure_figure.py

# 4. tests (unit + integration)
pytest -q # 74 passed, 1 skipped
```